A Survey of Quantization Methods for Efficient Neural Network Inference

Department of Computer Science & Artificial Intelligence

Abstract
This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction
Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding A Survey of Quantization Methods for Efficient Neural Network Inference.

Section topic: Multi-Paper Analysis
Question context: What conflicting findings exist regarding the effectiveness of post-training INT4 quantization on reasoning benchmarks?
Key Finding: Naive INT4 quantization severely degrades reasoning benchmarks [Survey of Quantization, p. 4], whereas activation-aware quantization protects critical weights and preserves accuracy [AWQ, p. 1].
Evidence passage: Standard round-to-nearest INT4 causes noticeable perplexity degradation on reasoning tasks
Activation-aware methods like AWQ retain near-lossless reasoning capability

2. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for A Survey of Quantization Methods for Efficient Neural Network Inference. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

3. Technical Methodology & Empirical Findings

Section topic: Optimization
Question context: What is Post-Training Quantization (PTQ) versus Quantization-Aware Training (QAT)?
Key Finding: PTQ quantizes weights and activations post-hoc without retraining, whereas QAT models quantization rounding noise during training for superior recovery of accuracy [Survey of Quantization Methods, p. 5].
Evidence passage: PTQ converts FP32 weights to INT8 using calibration data without re-training QAT simulates quantization errors in forward passes during training using fake quantization

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for A Survey of Quantization Methods for Efficient Neural Network Inference. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.