AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration

Department of Computer Science & Artificial Intelligence

Abstract
This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction

Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration.

Section topic: Optimization
Question context: How does AWQ (Activation-aware Weight Quantization) protect salient weights during INT4 compression?
Key Finding: AWQ identifies that weights associated with high-magnitude activation channels are disproportionately critical and applies channel-wise scaling to safeguard them from quantization distortion [AWQ, p. 3].
Evidence passage: Observes that only top 1% of weights corresponding to large activation channels are critical Scales salient weight channels prior to per-channel quantization to preserve output accuracy

Section topic: Multi-Paper Analysis
Question context: What conflicting findings exist regarding the effectiveness of post-training INT4 quantization on reasoning benchmarks?
Key Finding: Naive INT4 quantization severely degrades reasoning benchmarks [Survey of Quantization, p. 4], whereas activation-aware quantization protects critical weights and preserves accuracy [AWQ, p. 1].
Evidence passage: Standard round-to-nearest INT4 causes noticeable perplexity degradation on reasoning tasks Activation-aware methods like AWQ retain near-lossless reasoning capability

2. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

3. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.