Fast Inference from Large Language Models via Speculative Decoding

Department of Computer Science & Artificial Intelligence

Abstract
This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction
Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding Fast Inference from Large Language Models via Speculative Decoding.

Section topic: Optimization
Question context: What is speculative decoding and how does it achieve lossless latency reduction?
Key Finding: Speculative decoding uses a lightweight draft model to generate candidate token sequences which are then verified in parallel by the target model in a single execution step [Speculative Decoding, p. 3].
Evidence passage: A smaller draft model generates K speculative candidate tokens quickly The target model verifies all K tokens in parallel in a single forward pass

2. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for Fast Inference from Large Language Models via Speculative Decoding. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

3. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for Fast Inference from Large Language Models via Speculative Decoding. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for Fast Inference from Large Language Models via Speculative Decoding. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.