LoRA: Low-Rank Adaptation of Large Language Models

Department of Computer Science & Artificial Intelligence

Abstract
This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction
Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding LoRA: Low-Rank Adaptation of Large Language Models.

Section topic: Optimization
Question context: How does Low-Rank Adaptation (LoRA) freeze pre-trained weights while adapting LLMs?
Key Finding: LoRA freezes original weights W_0 and injects trainable rank-decomposition matrices delta_W = B * A with rank r << min(d, k), reducing trainable parameters by 10,000x [LoRA, p. 3].
Evidence passage: Decomposes weight update matrix into low-rank matrices: delta_W = B * A, where B in R^(d x r) and A in R^(r x k) with r << d Keeps pre-trained W_0 frozen and trains only low-rank matrices A and B

2. Technical Methodology & Empirical Findings

Section topic: Multi-Paper Analysis
Question context: Compare the parameter efficiency of prefix tuning, prompt tuning, and LoRA.
Key Finding: Prefix tuning inserts continuous prompt activations across all transformer layers, while LoRA modifies linear projection weights with low-rank decomposition matrices [LoRA, p. 2].
Evidence passage: Prefix tuning optimizes virtual tokens across all attention layers LoRA optimizes rank-decomposition matrices across attention projection weights

3. Technical Methodology & Empirical Findings

Section topic: Multi-Paper Analysis
Question context: Compare the memory savings achieved by LoRA versus QLoRA.
Key Finding: LoRA freezes 16-bit weights and introduces low-rank updates [LoRA, p. 3], whereas QLoRA compresses the base weights to 4-bit NF4, enabling 65B model fine-tuning on a single consumer GPU [QLoRA, p. 2].
Evidence passage: LoRA reduces trainable parameter memory by 10,000x but keeps base weights in 16-bit QLoRA compresses base weights to 4-bit NormalFloat, reducing base weight memory footprint by 4x

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for LoRA: Low-Rank Adaptation of Large Language Models. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.