SGDR: Stochastic Gradient Descent with Warm Restarts

Department of Computer Science & Artificial Intelligence

Abstract
This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction

Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding SGDR: Stochastic Gradient Descent with Warm Restarts.

2. Technical Methodology & Empirical Findings

Section topic: Optimization
Question context: How does cosine annealing with warm restarts modulate learning rates over training epochs?
Key Finding: Cosine annealing decreases the learning rate along a cosine curve before periodically resetting it to the maximum, enabling escape from sub-optimal local minima [SGDR, p. 4].
Evidence passage: Decays learning rate following a half-cosine curve and abruptly resets it to initial maximum Enables escaping sharp local minima and exploring multiple diverse basins

3. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for SGDR: Stochastic Gradient Descent with Warm Restarts. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for SGDR: Stochastic Gradient Descent with Warm Restarts. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.