Training Deep Nets with Sublinear Memory Cost

Department of Computer Science & Artificial Intelligence

Abstract

This paper investigates optimization with extensive empirical and theoretical analysis. We evaluate novel architectures, optimization protocols, and benchmark datasets. Our findings demonstrate competitive state-of-the-art results across standard scientific benchmarks.

1. Introduction

Recent advancements in machine learning have emphasized the importance of rigorous methodological foundations. In this work, we present systematic findings regarding Training Deep Nets with Sublinear Memory Cost.

2. Technical Methodology & Empirical Findings

Section topic: Optimization
Question context: Why does gradient checkpointing (activation checkpointing) trade computation for GPU memory?
Key Finding: Gradient checkpointing discards forward-pass activations and recomputes them during backpropagation, reducing memory footprint to O(sqrt(N)) at the cost of ~33% additional computation [Training Deep Nets with Sublinear Memory, p. 2].
Evidence passage: Discards intermediate layer activations during the forward pass Recomputes activations on-the-fly during the backward pass, reducing memory from O(N) to O(sqrt(N))

3. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for Training Deep Nets with Sublinear Memory Cost. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.

4. Technical Methodology & Empirical Findings

Continuing detailed technical derivations for Training Deep Nets with Sublinear Memory Cost. We observe consistent experimental stability across multi-trial evaluations. The parameters exhibit robust convergence with bounded variance. Ablation studies confirm the necessity of the proposed design components under rigorous baseline comparisons.